

Your reps go off-script. Your contract doesn't.

ClauseCatcher listens to live sales calls, catches any line that contradicts the signed contract, and reads the exact clause aloud while the customer is still on the line.

AssemblyAI Streaming STT v3

AssemblyAI Voice Agent

github.com/ishal1410/clausecatcher

LIVE CALL

Demo contract, 4 clauses

Rep

Seats are capped at fifty on this plan. ✓

Rep

We can also do a verbal discount for a big client.

Contradicts §3.1

Pricing & Seat Cap

"Flat \$48,000 up to 50 seats; extra seats need signed written amendment; no automatic or verbal discounting."

...

Reading the clause aloud...

A promise said out loud becomes the deal

REP SAID, ON THE CALL

“If you sign this week, we can do a verbal discount on the extra seats.”

SIGNED CONTRACT §3.1 SAYS

“Flat \$48,000 up to 50 seats; extra seats need signed written amendment; no automatic or verbal discounting.”

1

Nobody checks the call live

2

The gap surfaces after signing

3

Then it's a dispute, not a correction

Listen. Check. Speak the clause.

01

LISTEN

AssemblyAI Streaming STT v3

Always-on transcript of the rep, seeded with keyterms from the contract.

02

CHECK

Gemini claim-check

Each finished sentence is checked. The model returns a verdict and a clause ID, never text.

03

SPEAK

AssemblyAI Voice Agent

Reads the contract's literal clause aloud, and the cockpit shows the alert card.

Caught mid-call. Spoken back word for word.

- Rep offers a verbal discount
- §3.1 flagged in ~2 s
- Voice Agent reads §3.1 aloud
- Marked spoken verbatim



CALL STATUS
Off-contract

LINES CHECKED

4

CONTRADICTIONS

1

READ VERBATIM

0/1

CONTRADICTION ALERTS 1

 CONTRADICTION

§3.1 Pricing & Seat Cap

06:22:17 PM

REP SAID

"If you sign this week, we can do a verbal discount on the extra seats."

SIGNED CONTRACT §3.1 SAYS

"Flat \$48,000 up to 50 seats; extra seats need signed written amendment; no automatic or verbal discounting."

Reading the clause aloud...

CLAUSE WATCHLIST

4 monitored · 1 breached

 §3.1 Pricing & Seat Cap

Contradicted

§4.2 Renewal

Watching

§5.3 Data Retention

Watching

§6.1 Support SLA

Watching



Speaking to the call

AssemblyAI Voice Agent

ASK A CLAUSE ALOUD

Pick a section to read... Ctrl K

SIMULATE REP LINE

Demo aid: sent as if the rep said it.

We can do a verbal discount... 

Upload the contract. Every clause, quoted exactly.

§

Clauses extracted once, at upload

The words on screen are the words in the PDF.

K

Clause language seeds STT keyterms

Contract-specific terms transcribe more reliably.

✓

Consent confirmed before listening

The call can't start until disclosure is checked.

Your contract, clause by clause

These are the exact words every sentence on the call gets checked against.



demo-contract.pdf

4 clauses · read verbatim

Replace

4 clauses extracted

quoted exactly as written

§3.1 Pricing & Seat Cap

Flat \$48,000 up to 50 seats; extra seats need signed written amendment; no automatic or verbal discounting.

§4.2 Renewal

Auto-renews 12 months unless written notice 60 days before term ends; no mid-term termination.


§5.3 Data Retention

30-day recoverable archive then permanent deletion.

§6.1 Support SLA

P1 within 4 business hours Mon-Fri 9-6 ET; 24/7 requires Premium Support Addendum.

Two AssemblyAI connections, one job each



Why split? One Voice Agent session doing both jobs was tested live and started replying over the rep, even with a silence prompt. Streaming STT has no reply behavior, so silence is structural, not prompted. (ADR-0001)

We tested the easy paths. Both failed.

TRIED AND DROPPED

One Voice Agent session for everything

Auto-started 4 replies despite a silence prompt

Alert injected as a fake user message

Agent ignored the injected content

SHIPPED

Streaming STT v3 listens all call

Pure transcription: nothing to talk over

Voice Agent reply.create: “say exactly”

Spoken alert matched the clause, similarity 1.0

Both products are load-bearing. Cut Streaming STT and there is no transcript to check. Cut the Voice Agent and the alert is a silent card instead of an interruption the rep hears.

It never puts words in a rep's mouth



Literal clause text only

Everything spoken or shown comes from the uploaded contract. The LLM only picks which clause applies.



Errors ≠ accusations

A failed, timed-out or unparseable check resolves to “unclear”, never a contradiction.



Evidence, or no alert

Every alert cites a specific clause ID. Each clause fires at most once per 20 seconds.

A false accusation is worse than a missed one.

One live end-to-end run, real connections

~2 s

sentence end to alert

1.0

spoken vs. contract
similarity

~78 ms

speak request to first
audio

~\$0.12

API usage, full demo
call

- False claims on §3.1 and §6.1 both caught; consistent lines raised nothing
- A monitor's spoken question about §4.2 answered correctly by voice
- 136 server tests passing

Run of 2026-09-15: real AssemblyAI Streaming STT, Voice Agent and Gemini connections, scripted rep lines fed through the pipeline.

BUILT FOR

- Sales-ops teams running live calls
- Compliance teams who own the contract
- ~\$0.12 API usage per demo call

NEXT

- Browser-mic path, hardened end to end
- A hosted, deployed instance
- Multiple contracts per call

The correction a rep hears is the contract itself, spoken while the call is still happening.